

# SQL/JSON

---

SDRP Lecture 6  
Christian Thomsen

[The slides are largely based on the [PostgreSQL documentation](#)]

# Agenda

---

- JSON basics
- Storing and querying JSON in PostgreSQL
- SQL/JSON path expressions
- From tables to JSON

# JSON

---

- JSON = JavaScript Object Notation
- Open standard for data exchange
- Language independent
- Uses text only, UTF-8 encoded
- Low overhead
  - Shorter/less verbose than XML
    - ◆ XML: `<SomeTag><AnotherTag>Data</AnotherTag></SomeTag>`
  - Easier to parse
- Like XML, JSON
  - is self-describing/human readable
  - supports nesting
  - is widely supported in different programming languages

# JSON basics

- A value can be
  - A number – integer or floating point, **but not** NaN or Inf
  - A string – in quotes
  - A Boolean – **true** or **false**
  - **null**
  - An ordered array – [**e1**, **e2**, **e3**, **e4**], possibly of different types
  - An object of unordered key/value pairs (where the keys are strings):  
`{"k1":v1, "k2":v2, "k3":v3}`

Example of document:

```
• { "name"      : "Rishi Sunak",  
    "address"  : { "street" : "Downing Street",  
                  "number"  : 10,  
                  "city"    : "London"  
                },  
    "age"      : 42  
  }
```

# Storing JSON in your relational database

- Microsoft gives the following reasons:
  - Simplify complex models
  - Store retail and e-commerce data
  - Process log and telemetry data
  - Store semi-structured IoT data
  - Simplify REST API data
- ISO/IEC argue in TR 19075-6:2017(E)
  - “technical, business, and government worlds are increasingly using both” [relational and JSON data]
  - “There are great benefits when a single data management system can concurrently handle all of the data”
    - ◆ reduced administrative costs
    - ◆ improved security and transactions management
    - ◆ better performance and greater optimizability
    - ◆ better resource allocation
    - ◆ increased developer productivity

# JSON in PostgreSQL

---

- PostgreSQL has good support for JSON
- Two types for JSON storage
  - **json**: Stores an exact copy of the input (preserves whitespace and order)
  - **jsonb**: Stores in binary format (does *not* preserve whitespace and order)
    - ◆ Slightly slower to input, much faster to process
    - ◆ Indexable
  - We use jsonb in this lecture

# Example

```
create table courses(id int, cnt jsonb);  
-- using the new jsonb data type  
  
insert into courses values  
(1, '{"id":4, "name":"OOP", "semester":2}'),  
(2, '{"id":2, "name":"SDRP", "semester":7}');  
-- a plain insert statement  
-- Note that the id in the JSON is different  
-- from the value in the column id!  
  
insert into courses values(3, '{"invalid"}');  
-- Fails. Good! Would not happen with TEXT
```

# Operators

- You can extract an object field or element with ->
- `select cnt->'name' as name,  
cnt->'semester' as sem  
from courses;`
- For x->n
  - n must be a key name (text) when x is an object
  - n must be an integer when x is an array
- -> gives back jsonb so you can chain ->
  - `select ('{"a": [9,8,7,6]} '::jsonb) ->'a' ->0;`  
gives the result **9** as jsonb
- If you want text back, you can use ->>
  - `select ('{"a": [9,8,7,6]} '::jsonb) ->'a' ->>0;`  
gives the result **9** as text

name	sem
OOP	2
SDRP	7



# Operators

- With `#>` you can specify a *path* of keys and array indices to the element to extract
- `select ('{"a": {"b": ["foo", "bar"]}}' :: jsonb) #>' {a, b, 1} ' ;` gives **"bar"** (as jsonb)
- `#>>` returns text
- `select '{"a": {"b": ["foo", "bar"]}}' :: jsonb  
->'no' ->'such' ->'keys' ;`  
and  
`select '{"a": {"b": ["foo", "bar"]}}' :: jsonb  
#>' {no, such, keys} ' ;`  
both return **null** – no error

# Existence

- You can check if a string exists as a top-level key or array element with ?
- `select '{"a":1, "b":2}'::jsonb ? 'b';`  
gives **true**
- `select '{"a": ["hello", "world"]}'::jsonb  
->'a' ? 'world';`  
also gives **true**
- Find the courses where there is info about the TA  
`select cnt from courses where cnt ? 'TA';`

# Containment

- You can check if an object is *contained* in another with @>
- The contained object must match the containing object wrt. structure and data contents
  - The containing object may contain more
  - Order of array elements does not matter and duplicates do not matter
- Examples
  - `select '[1, 2, 3]':::jsonb @> '[1, 3]':::jsonb; is true`
  - `select '[1, 2, 3]':::jsonb @> '[3, 1]':::jsonb; is true`
  - `select '[1, 2, 3]':::jsonb @> '[1, 2, 2]':::jsonb; is true`

# Containment, cont.

- `select '{"licensePlate": "ABC123", "brand": "Volvo", "owner": {"name": "Bo Xi", "phone": 123}}'::jsonb @> '{"brand": "Volvo"}' is true`
- `select '{"licensePlate": "ABC123", "brand": "Volvo", "owner": {"name": "Bo Xi", "phone": 123}}'::jsonb @> '{"phone": 123}' is false`
- `select '{"licensePlate": "ABC123", "brand": "Volvo", "owner": {"name": "Bo Xi", "phone": 123}}'::jsonb @> '{"owner": {"phone": 123}}' is true`
- Often, you will use @> in the WHERE clause
- Find all 7th semester courses:
  - `select cnt from courses where cnt @> '{"semester": 7}' ;`

# Querying with subscripting

- Subscripting can be used to extract and modify elements

- `select cnt['name'] as name,  
cnt['semester'] as sem  
from courses;`

name	sem
OOP	2
SDRP	7

- `select ('[0,1,2]'::jsonb)[0] gives 0`
- `select ('[0,1,2]'::jsonb)[-1] gives 2`
- `select ('[0,1,2]'::jsonb)[3] gives [null]`
- `select ('{"a": {"b": {"c": 1}}}'::jsonb)  
['a']['b']['c'] gives 1`
- `select ('{"a": {"b": {"c": 1}}}'::jsonb)  
['a']['b']['c']['d'] gives [null]`

# Updating with subscripting

- If `val['a']` or `val['a']['b']` is not defined, an empty object will be created and filled as necessary
  - The last *existing* element must be an object or array
- **update courses**  
`set cnt['TA']['prereqs'] = 'SDRP'`  
**where id=2**  
will create an object for “TA” and then add the k/v “prereqs”:“SDRP” to that

# Agenda

---

- JSON basics
- Storing and querying JSON in PostgreSQL
- **SQL/JSON path expressions**
- From tables to JSON

# SQL/JSON path expressions

---

- PostgreSQL also has the type **jsonpath** for efficient binary representation of SQL/JSON path expressions
  - You input the path expressions as SQL strings, i.e., in single quotes
- SQL/JSON path expressions are used to specify items to be retrieved
  - Similar to XPath expressions to retrieve from XML



# SQL/JSON path expressions

- A path expression is a sequence of path elements:
  - Literals: Text, numbers, true/false, null
  - Variables:
    - ◆ \$ is the *context item* and represents the JSON value being queried
    - ◆ \$name is a named variable
    - ◆ @ represents a value being filtered
  - Accessors:
    - ◆ .key gives the object member with that key
    - ◆ .\* gives the values for all members at the top level of the current obj.
    - ◆ [idx] gives the single element at that position in an array  
(0-based – **not** 1-based as regular SQL arrays, **[last]** also legal)
    - ◆ [start **to** end] gives an array slice
    - ◆ [\*] gives all array elements
- Evaluated left to right - an operator deals with the result of the previous step. Parentheses can be added

# Example

- `select jsonb_path_query(cnt, '$.name') from courses;`

"OOP"
"SDRP"

- `select jsonb_path_query('{"a": [0,1,2,3]}',  
'$.a[0 to 1]')`

0
1

# Jsonpath operators and methods

- The "usual" operators: + - \* / %
  - Only unary + and – can iterate over multiple values
  - `select jsonb_path_query('{ "a": [0,1,2,3] }', '$.a[0] + 1')` works
  - `select jsonb_path_query('{ "a": [0,1,2,3] }', '$.a + 1')` **fails**
- `.size()` → size of array or 1 if not on an array
- `.type()` → "string", "object", ...
- `.double()`
- `.ceiling()`
- `.floor()`
- `.abs()`
- `.datetime()` – converts a string to date/time value
  - Order of attempts: date, timetz, time, timestamptz, timestamp

# Filters

- Filter expressions can be used – similar to SQL WHERE
- ? (condition)
- Written immediately after the step to apply the filter to
  - @ denotes the value being filtered coming from the prev. step
- `select jsonb_path_query(cnt, '$ ? (@.semester >= 6).name') from courses;`  
gives “**SDRP**”
- You can of course have more filters:
  - \$ ? (@.x < 10) ? (@.y < 20)
  - \$ ? (@.x < 10).y ? (@ < 20)
  - \$ ? (@.x < 10 && @.y < 20)

## Filters, cont.

- The expected operators can be used: ==, !=, <>, <, >, <=, >=, &&, ||, !

In addition

- **is unknown:**

```
jsonb_path_query('[-1, 2, 7, "foo"]',  
'$[*] ? (@ > 0) is unknown') gives "foo"
```

- **starts with:**

```
jsonb_path_query('["John Smith", "Ann Ai"]',  
'$[*] ? (@ starts with "John")')  
gives "John Smith"
```

- **exists:**

```
select jsonb_path_query(cnt, '$ ? (exists  
(@.TA)).name') from courses; gives "SDRP"
```

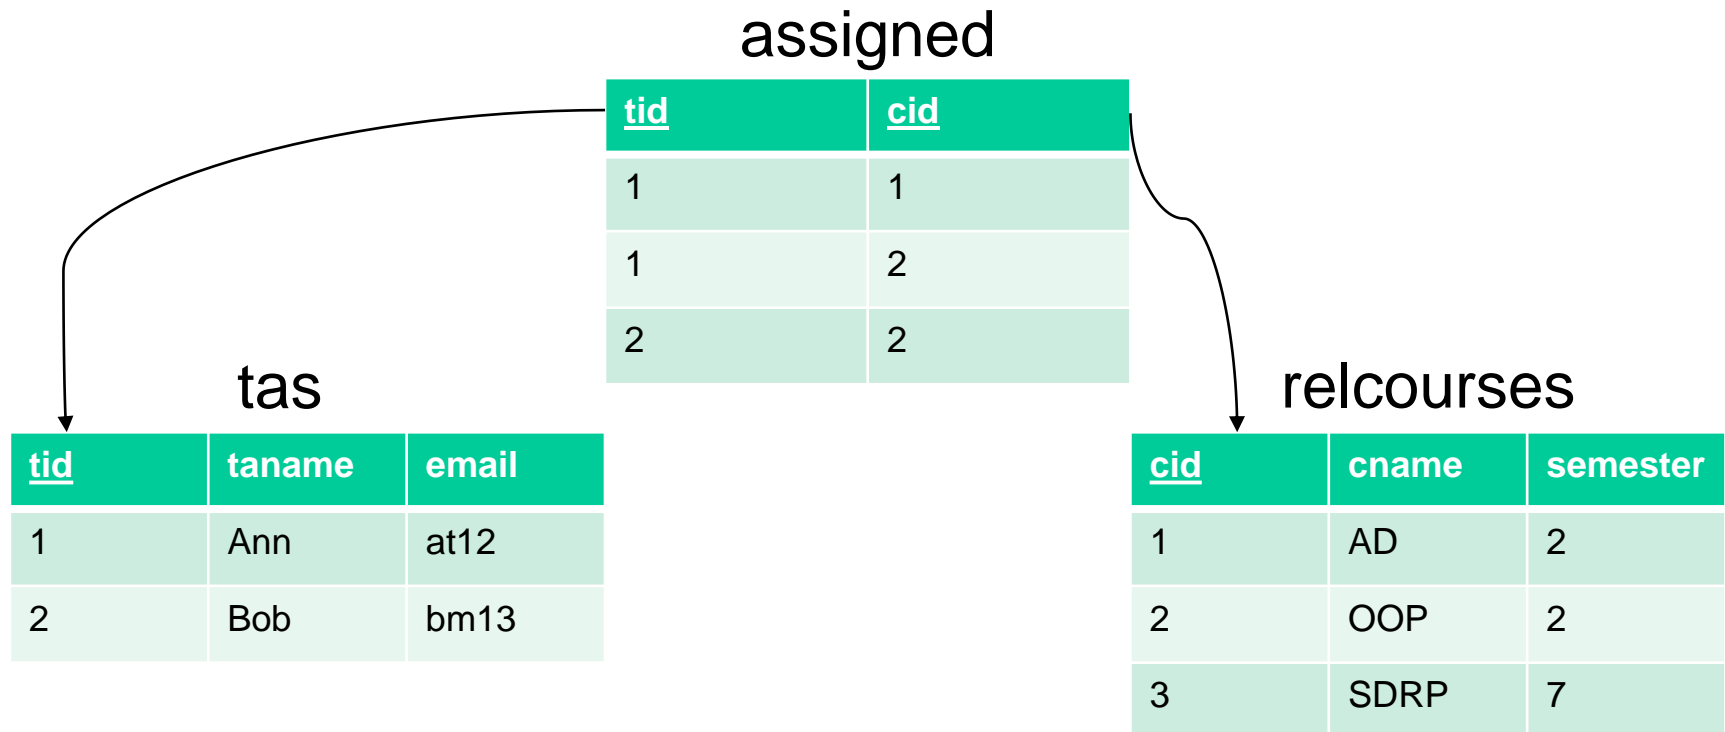
# Agenda

---

- JSON basics
- Storing and querying JSON in PostgreSQL
- SQL/JSON path expressions
- **From tables to JSON**

# From tables to JSON

- Much data is stored in relational databases, but you may need to exchange that with someone else
- You can create JSON from your relational data
- Setting the stage for our example:



# Simple start

- `select to_jsonb(relcourses.*)  
from relcourses;`

`to_jsonb`

`"{"cid": 1, "cname": "AD", "semester": 2}"`

`"{"cid": 2, "cname": "OOP", "semester": 2}"`

`"{"cid": 3, "cname": "SDRP", "semester": 7}"`

- So far so good, but we lack information about the TAs
- `select jsonb_build_object('name', taname,  
'email', email) from tas;`

`jsonb_build_object`

`"{"name": "Ann", "email": "at12"}"`

`"{"name": "Bob", "email": "bm13"}"`



# Courses and TAs

- with  
tainfo as (select tid, jsonb\_build\_object('name',  
  taname, 'email', email) as ta from tas),  
coursesandtas as (select cname, semester, ta  
  from relcourses natural left outer join  
  assigned natural left outer join tainfo)  
select to\_jsonb(coursesandtas.\*) from coursesandtas;

## to\_jsonb

"{"ta": {"name": "Ann", "email": "at12"}, "cname": "AD", "semester": 2}"

"{"ta": {"name": "Ann", "email": "at12"}, "cname": "OOP", "semester": 2}"

"{"ta": {"name": "Bob", "email": "bm13"}, "cname": "OOP", "semester": 2}"

"{"ta": null, "cname": "SDRP", "semester": 7}"

OOP is represented twice  
because it has two TAs

# Courses and TAs

- with  
tainfo as (select tid, jsonb\_build\_object('name',  
taname, 'email', email) as ta from tas),  
coursesandtas as  
(select cname, semester, jsonb\_agg(ta) as ta  
from relcourses natural left outer join  
assigned natural left outer join tainfo  
group by cname, semester)  
select to\_jsonb(coursesandtas.\*) from coursesandtas;

to\_jsonb

"{"ta": [null], "cname": "SDRP", "semester": 7}"

"{"ta": [{"name": "Ann", "email": "at12"}, {"name": "Bob", "email": "bm13"}],  
"cname": "OOP", "semester": 2}"

"{"ta": [{"name": "Ann", "email": "at12"}], "cname": "AD", "semester": 2}"

OOP is now represented once! But we  
want a single object for all courses

# Making a single document

- with

```
tainfo as (select tid, jsonb_build_object('name',
      taname, 'email', email) as ta from tas),
coursesandtas as
      (select cname, semester, jsonb_agg(ta) as ta
      from relcourses natural left outer join
      assigned natural left outer join tainfo
      group by cname, semester),
jsonparts as (select to_jsonb(coursesandtas.*) as json
      from coursesandtas),
jsonarray as (select jsonb_agg(json) as jsonarray from
      jsonparts)
select jsonb_build_object('courses', jsonarray)
from jsonarray;
```

# Getting "pretty" JSON

```
select jsonb_pretty(jsonb_build_object(...as before...))
```

jsonb\_pretty

```
"{
  "courses": [
    {
      "ta": [
        null
      ],
      "cname": "SDRP",
      "semester": 7
    },
    {
      "ta": [
        {
          "name": "Ann",
          "email": "at12"
        },
        {
          "name": "Bob",
          "email": "bm13"
        }
      ],
      "cname": "OOP",
      "semester": 2
    },
    {
      "ta": [
        {
          "name": "Ann",
          "email": "at12"
        }
      ],
      "cname": "AD",
      "semester": 2
    }
  ]
}
```

"No TA" should not be represented by "[null]". We fix that in the exercises

# Summary

---

- JSON basics
- Storing and querying JSON in PostgreSQL
- SQL/JSON path expressions
- From tables to JSON
- Exercises on Moodle